

Metric Learning

ICML 2010 Tutorial

Brian Kulis
University of California at Berkeley

June 21, 2010

Introduction

Learning problems with distances and similarities

- k -means
- Support vector machines
- k -nearest neighbors
- Most algorithms that employ kernel methods
- Other clustering algorithms (agglomerative, spectral, etc)
- ...

Choosing a distance function




Choosing a distance function


Example: UCI Wine data set

- 13 features
 - 9/13 features have mean value in $[0, 10]$
 - 3/13 features have mean value in $[10, 100]$
 - One feature has a mean value of 747 (with std 315)
- Using a standard distance such as Euclidean distance, the largest feature dominates the computation
 - That feature may not be important for classification
- Need a weighting of the features that improves classification or other tasks

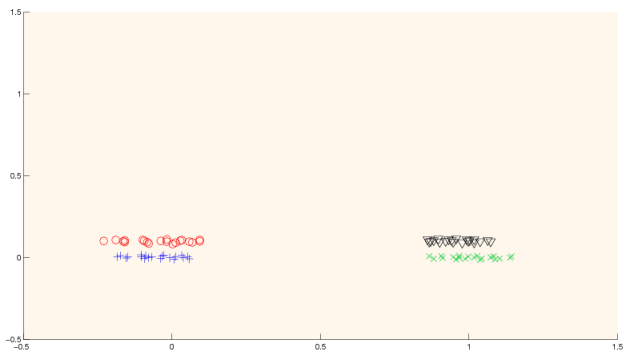
Example: Four Blobs




Example: Four Blobs



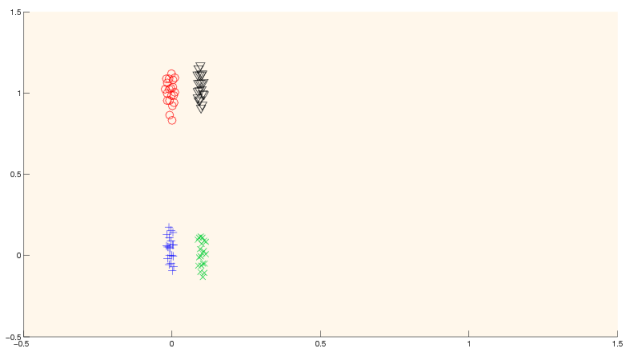
Example: Four Blobs



Example: Four Blobs



Example: Four Blobs



Metric Learning as Learning Transformations

- Feature re-weighting
 - Learn weightings over the features, then use standard distance (e.g., Euclidean) after re-weighting
 - Diagonal Mahalanobis methods (e.g., Schultz and Joachims)
 - Number of parameters grows linearly with the dimensionality d
- Full linear transformation
 - In addition to scaling of features, also rotates the data
 - For transformations from d dimensions to d dimensions, number of parameters grows quadratically in d
 - For transformations to $r < d$ dimensions, this is linear dimensionality reduction
- Non-linear transformation
 - Variety of methods
 - Neural nets
 - Kernelization of linear transformations
 - Complexity varies from method to method

Supervised vs Unsupervised Metric Learning

- Unsupervised Metric Learning
 - Dimensionality reduction techniques
 - Principal Components Analysis
 - Kernel PCA
 - Multidimensional Scaling
 - In general, not the topic of this tutorial...
- Supervised and Semi-supervised Metric Learning
 - Constraints or labels given to the algorithm
 - Example: set of similarity and dissimilarity constraints
 - This is the focus of the tutorial

Themes of the tutorial

- Not just a list of algorithms
 - General principles
 - Focus on a few key methods
- Recurring ideas
 - Scalability
 - Linear vs non-linear
 - Online vs offline
 - Optimization techniques utilized
 - Statements about general formulations
- Applications
 - Where is metric learning applied?
 - Success stories
 - Limitations

Outline of Tutorial

- Motivation
- Linear metric learning methods
 - Mahalanobis metric learning
 - Per-example methods
- Non-linear metric learning methods
 - Kernelization of Mahalanobis methods
 - Other non-linear methods
- Applications
- Conclusions

ON THE GENERALIZED DISTANCE IN STATISTICS.

By P. C. MAHALANOBIS.

(Read January 4, 1936.)

1. A normal (Gauss-Laplacian) statistical population in P -variates is usually described by a P -dimensional frequency distribution :—

$$df = \text{const.} \times e^{-\frac{1}{2\alpha} [A_{11}(x_1 - \alpha_1)^2 + A_{22}(x_2 - \alpha_2)^2 + \dots + 2A_{12}(x_1 - \alpha_1)(x_2 - \alpha_2) + \dots]} \cdot dx_1 \cdot dx_2 \dots dx_P \quad (1.0)$$

where

$\alpha_1, \alpha_2, \dots, \alpha_P$ = the population (mean) values
of the P -variates x_1, x_2, \dots, x_P (1.1)

$\alpha_{ii} = \sigma_i^2$, are the respective variances (1.2)

$\alpha_{ij} = \sigma_i \cdot \sigma_j \cdot \rho_{ij}$, where ρ_{ij} = the coefficient of correlation between the i th and j th variates (1.3)

α is the determinant $|\alpha_{ij}|$ defined more fully in (2.2), and A_{ij} is the minor of α_{ij} in this determinant.

A P -variate normal population is thus completely specified by the set of $P(P+1)/2$ parameters* :—



Mahalanobis Distances

- Assume the data is represented as N vectors of length d :

$$X = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N]$$

- Squared Euclidean distance



$$\begin{aligned}d(\mathbf{x}_1, \mathbf{x}_2) &= \|\mathbf{x}_1 - \mathbf{x}_2\|_2^2 \\ &= (\mathbf{x}_1 - \mathbf{x}_2)^T (\mathbf{x}_1 - \mathbf{x}_2)\end{aligned}$$

- Let $\Sigma = \sum_{i,j} (\mathbf{x}_i - \mu)(\mathbf{x}_j - \mu)^T$
- The “original” Mahalanobis distance:

$$d_M(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T \Sigma^{-1} (\mathbf{x}_1 - \mathbf{x}_2)$$

Mahalanobis Distances

- Equivalent to applying a *whitening transform*



Mahalanobis Distances

- Assume the data is represented as N vectors of length d :

$$X = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N]$$

- Squared Euclidean distance

$$\begin{aligned}d(\mathbf{x}_1, \mathbf{x}_2) &= \|\mathbf{x}_1 - \mathbf{x}_2\|_2^2 \\ &= (\mathbf{x}_1 - \mathbf{x}_2)^T (\mathbf{x}_1 - \mathbf{x}_2)\end{aligned}$$

- Mahalanobis distances for metric learning
 - Distance parametrized by $d \times d$ positive semi-definite matrix A :

$$d_A(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)$$

- Used for many existing metric learning algorithms

[Xing, Ng, Jordan, and Russell; NIPS 2002]

[Bar-Hillel, Hertz, Shental, and Weinshall; ICML 2003]

[Bilenko, Basu, and Mooney; ICML 2004]

[Globerson and Roweis; NIPS 2005]

[Weinberger, Blitzer, and Saul; NIPS 2006]

Mahalanobis Distances

$$d_A(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)$$

- Why is A positive semi-definite (PSD)?
 - If A is not PSD, then d_A could be negative
 - Suppose $\mathbf{v} = \mathbf{x}_1 - \mathbf{x}_2$ is an eigenvector corresponding to a negative eigenvalue λ of A

$$\begin{aligned}d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\ &= \mathbf{v}^T A \mathbf{v} \\ &= \lambda \mathbf{v}^T \mathbf{v} = \lambda < 0\end{aligned}$$

Mahalanobis Distances

- Properties of a metric:
 - $d(\mathbf{x}, \mathbf{y}) \geq 0$
 - $d(\mathbf{x}, \mathbf{y}) = 0$ if and only if $\mathbf{x} = \mathbf{y}$
 - $d(\mathbf{x}, \mathbf{y}) = d(\mathbf{y}, \mathbf{x})$
 - $d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z})$
- d_A is not technically a metric
 - Analogous to Euclidean distance, need the square root:

$$\sqrt{d_A(\mathbf{x}_1, \mathbf{x}_2)} = \sqrt{(\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2)}$$

- Square root of the Mahalanobis distance satisfies all properties if A is strictly positive definite, but if A is positive semi-definite then second property is not satisfied
 - Called a *pseudo-metric*
- In practice, most algorithms work only with d_A


Mahalanobis Distances

- Can view d_A as the squared Euclidean distance after applying a linear transformation
 - Decompose $A = G^T G$ via Cholesky decomposition
 - (Alternatively, take eigenvector decomposition $A = V\Lambda V^T$ and look at $A = (\Lambda^{1/2} V^T)^T (\Lambda^{1/2} V^T)$)
- Then we have


$$\begin{aligned}d_A(\mathbf{x}_1, \mathbf{x}_2) &= (\mathbf{x}_1 - \mathbf{x}_2)^T A (\mathbf{x}_1 - \mathbf{x}_2) \\ &= (\mathbf{x}_1 - \mathbf{x}_2)^T G^T G (\mathbf{x}_1 - \mathbf{x}_2) \\ &= (G\mathbf{x}_1 - G\mathbf{x}_2)^T (G\mathbf{x}_1 - G\mathbf{x}_2) \\ &= \|G\mathbf{x}_1 - G\mathbf{x}_2\|_2^2\end{aligned}$$

- Mahalanobis distance is just the squared Euclidean distance after applying the linear transformation G


Example: Four Blobs



Example: Four Blobs




Example: Four Blobs




- Want to learn:

$$A = \begin{pmatrix} 1 & 0 \\ 0 & \epsilon \end{pmatrix} \quad G = \begin{pmatrix} 1 & 0 \\ 0 & \sqrt{\epsilon} \end{pmatrix}$$

Example: Four Blobs



Example: Four Blobs




- Want to learn:

$$A = \begin{pmatrix} \epsilon & 0 \\ 0 & 1 \end{pmatrix} \quad G = \begin{pmatrix} \sqrt{\epsilon} & 0 \\ 0 & 1 \end{pmatrix}$$

Drawbacks to Mahalanobis Metric Learning

- Memory overhead grows quadratically with the dimensionality of the data
 - Does not scale to high-dimensional data ($d = O(10^6)$ for many image embeddings)
- Only works for linearly separable data



- **Cannot seemingly be applied to “real” data!**
- These drawbacks will be discussed later

Metric Learning Problem Formulation

- Typically 2 main pieces to a Mahalanobis metric learning problem
 - A set of constraints on the distance
 - A regularizer on the distance / objective function
- In the constrained case, a general problem may look like:

$$\begin{aligned} \min_A \quad & r(A) \\ \text{s.t.} \quad & c_i(A) \leq 0 \quad 0 \leq i \leq C \\ & A \succeq 0 \end{aligned}$$

- r is a regularizer/objective on A and c_i are the constraints on A
- An unconstrained version may look like:

$$\min_{A \succeq 0} r(A) + \lambda \sum_{i=1}^C c_i(A)$$

Defining Constraints

- Similarity / Dissimilarity constraints

- Given a set of pairs \mathcal{S} of points that should be similar, and a set of pairs of points \mathcal{D} of points that should be dissimilar
- A single constraint would be of the form

$$d_A(\mathbf{x}_i, \mathbf{x}_j) \leq \ell$$

for $(i, j) \in \mathcal{S}$ or

$$d_A(\mathbf{x}_i, \mathbf{x}_j) \geq u$$

for $(i, j) \in \mathcal{D}$

- Easy to specify given class labels
- Relative distance constraints
- Given a triple $(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k)$ such that the distance between \mathbf{x}_i and \mathbf{x}_j should be smaller than the distance between \mathbf{x}_i and \mathbf{x}_k , a single constraint is of the form

$$d_A(\mathbf{x}_i, \mathbf{x}_j) \leq d_A(\mathbf{x}_i, \mathbf{x}_k) - m,$$

where m is the margin

- Popular for ranking problems

Defining Constraints

- Aggregate distance constraints
 - Constrain the sum of all pairs of same-class distances to be small, e.g.,

$$\sum_{ij} y_{ij} d_A(\mathbf{x}_i, \mathbf{x}_j) \leq 1$$

where $y_{ij} = 1$ if \mathbf{x}_i and \mathbf{x}_j are in the same class, and 0 otherwise

- Other constraints
 - Non-parametric probability estimation constraints
 - Constraints on the generalized inner product $\mathbf{x}_i^T \mathbf{A} \mathbf{x}_j$:

$$d_A(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{A} \mathbf{x}_i + \mathbf{x}_j^T \mathbf{A} \mathbf{x}_j - 2\mathbf{x}_i^T \mathbf{A} \mathbf{x}_j$$

Defining the Regularizer or Objective

- Loss/divergence functions
 - Squared Frobenius norm: $\|A - A_0\|_F^2$
 - LogDet divergence: $\text{tr}(AA_0^{-1}) - \log \det(AA_0^{-1}) - d$
 - General loss functions $D(A, A_0)$
 - Will discuss several of these later
- Other regularizers
 - $\|A\|_F^2$
 - $\text{tr}(AC_0)$ (i.e., if C_0 is the identity, this is the trace norm)

Choosing a Regularizer

- Depends on the problem!
- Example 1: $\text{tr}(A)$
 - Trace function is the sum of the eigenvalues
 - Analogous to the ℓ_1 penalty, promotes sparsity
 - Leads to low-rank A
- Example 2: LogDet Divergence
 - Defined only over positive semi-definite matrices
 - Makes computation simpler
 - Possesses other desirable properties
- Example 3: $\|A\|_F^2$
 - Arises in many formulations
 - Easy to analyze and optimize

Defining the Optimization

- Many existing Mahalanobis distance learning methods can be obtained simply by choosing a regularizer/objective and constraints
- We will discuss properties of several of these

Xing et al.'s MMC

Problem posed as follows:

$$\begin{aligned} \max_A \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{D}} \sqrt{d_A(\mathbf{x}_i, \mathbf{x}_j)} \\ \text{s.t.} \quad & c(A) = \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) \leq 1 \\ & A \succeq 0. \end{aligned}$$

- Here, \mathcal{D} is a set of pairs of dissimilar pairs, \mathcal{S} is a set of similar pairs
- Objective tries to maximize sum of dissimilar distances
- Constraint keeps sum of similar distances small
 - Use square root in regularizer to avoid trivial solution

[Xing, Ng, Jordan, and Russell; NIPS 2002]

Algorithm

- Based on gradient descent over the objective followed by an iterative projection step to find a feasible A
 - Constraint $c(A)$ is linear in A , can be solved cheaply
 - Orthogonal projection onto $A \succeq 0$ achieved by setting A 's negative eigenvalues to 0
 - Iterative between these two steps to find feasible A for both constraints, then take a step in the gradient of the objective
- Despite relative simplicity, the algorithm is fairly slow (many eigenvalue decompositions required)
- Does not scale to large problems
- Objective and constraints only look at the sums of distances

Problem formulated as follows:

$$\begin{aligned} \min_A \quad & \|A\|_F^2 \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (i, j, k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

- Constraints in \mathcal{R} are relative distance constraints
- There may be no solution to this problem; introduce slack variables

$$\begin{aligned} \min_{A, \xi} \quad & \|A\|_F^2 + \gamma \sum_{(i, j, k) \in \mathcal{R}} \xi_{ijk} \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (i, j, k) \in \mathcal{R} \\ & \xi_{ijk} \geq 0 \quad \forall (i, j, k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

Algorithm

- Key simplifying assumption made
 - $A = M^T D M$, where M is assumed fixed and known and D is diagonal

$$\begin{aligned}d_A(\mathbf{x}_i, \mathbf{x}_j) &= (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \\ &= (\mathbf{x}_i - \mathbf{x}_j)^T M^T D M (\mathbf{x}_i - \mathbf{x}_j) \\ &= (M\mathbf{x}_i - M\mathbf{x}_j)^T D (M\mathbf{x}_i - M\mathbf{x}_j)\end{aligned}$$

- Effectively constraining the optimization to diagonal matrices
- Resulting optimization problem is very similar to SVMs, and resulting algorithm is similar
- By choosing M to be a matrix of data points, method can be kernelized
- Fast algorithm, but less general than full Mahalanobis methods

Problem formulated as follows:

$$\begin{aligned} \min_{A, \xi, \gamma} \quad & \|A\|_F^2 + \frac{C_S}{N_S} \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in S} \xi_{ij} + \frac{C_D}{N_D} \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in D} \xi_{ij} - C_D \gamma \nu \\ \text{s.t.} \quad & d_l(\mathbf{x}_i, \mathbf{x}_j) \geq d_A(\mathbf{x}_i, \mathbf{x}_j) - \xi_{ij} \quad \forall (\mathbf{x}_i, \mathbf{x}_j) \in S \\ & d_A(\mathbf{x}_i, \mathbf{x}_j) - d_l(\mathbf{x}_i, \mathbf{x}_j) \geq \gamma - \xi_{ij} \quad \forall (\mathbf{x}_i, \mathbf{x}_j) \in D \\ & \xi_{ij} \geq 0 \\ & \gamma \geq 0 \\ & A \succeq 0. \end{aligned}$$

- Same regularization as Schultz and Joachims
- Similarity/dissimilarity constraints instead of relative distance constraints
- No simplifying assumptions made about A

[Kwok and Tsang; ICML 2003]

Neighbourhood Components Analysis

Problem formulated as follows:

$$\begin{aligned} \max_A \quad & \sum_i \sum_{j \in C_i, j \neq i} \frac{\exp(-d_A(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k \neq i} \exp(-d_A(\mathbf{x}_i, \mathbf{x}_k))} \\ \text{s.t.} \quad & A \succeq 0. \end{aligned}$$

- C_i is the set of points in the same class as point \mathbf{x}_i (not including \mathbf{x}_i)

Motivation

- Minimize the leave-one-out KNN classification error
 - LOO error function is discontinuous
 - Replace by a softmax; each point \mathbf{x}_i chooses a nearest neighbor \mathbf{x}_j based on probability

$$p_{ij} = \frac{\exp(-d_A(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k \neq i} \exp(-d_A(\mathbf{x}_i, \mathbf{x}_k))}$$

[Goldberger, Roweis, Hinton, and Salakhutdinov; NIPS 2004]

Neighbourhood Components Analysis

Algorithm

- Problem is *non-convex*
- Rewrite in terms of G , where $A = G^T G$
 - Eliminates $A \succeq 0$ constraint
- Run gradient descent over G

Properties

- Easy to control the rank of A : just optimize over low-rank G
- Simple, unconstrained optimization
- No guarantee of global solution

Recall NCA probabilities

$$p_{ij} = \frac{\exp(-d_A(\mathbf{x}_i, \mathbf{x}_j))}{\sum_{k \neq i} \exp(-d_A(\mathbf{x}_i, \mathbf{x}_k))}$$

- Introduce an “ideal” probability distribution p_{ij}^0 :

$$p_{ij}^0 \propto \begin{cases} 1 & \text{if } i \text{ and } j \text{ from same class} \\ 0 & \text{otherwise.} \end{cases}$$

- Minimize divergence between p^0 and p :

$$\begin{aligned} \min_A \quad & KL(p^0, p) \\ \text{s.t.} \quad & A \succeq 0. \end{aligned}$$

[Globerson and Roweis; NIPS 2005]

Properties

- Unlike NCA, MCML is *convex*
- Global optimization possible
 - Algorithm based on optimization over the dual
 - Similar to Xing: gradient step plus projection
 - Not discussed in detail in this tutorial

A Closer Look at Some Algorithms

- As can be seen, several objectives are possible
- We will take a look at 3 algorithms in-depth for their properties
- POLA
 - An *online* algorithm for learning metrics with provable regret
 - Also the first supervised metric learning algorithm that was shown to be kernelizable
- LMNN
 - Very popular method
 - Algorithm scalable to billions of constraints
 - Extensions for learning multiple metrics
- ITML
 - Objective with several desirable properties
 - Simpler kernelization construction
 - Online variant

Setup

- Estimate d_A in an online manner
- Also estimate a threshold b
- At each step t , observe tuple $(\mathbf{x}_t, \mathbf{x}'_t, y_t)$
 - $y_t = 1$ if \mathbf{x}_t and \mathbf{x}'_t should be similar; -1 otherwise
- Consider the following loss

$$\ell_t(A, b) = \max(0, y_t(d_A(\mathbf{x}_t, \mathbf{x}'_t)^2 - b) + 1)$$

- Hinge loss, margin interpretation
- Appropriately update A and b each iteration

[Shalev-Shwartz, Singer, and Ng; NIPS 2004]

Online Learning Setup

- Define the total loss to be

$$L = \sum_{t=1}^T \ell_t(A_t, b_t)$$

- Online learning methods compare the loss to the best *fixed*, offline predictor A^* and threshold b^*
- Define *regret* for T total timesteps as

$$R_T = \sum_{t=1}^T \ell_t(A_t, b_t) - \sum_{t=1}^T \ell_t(A^*, b^*)$$

- Design an algorithm that minimizes the regret
 - Same setup as in other online algorithms (classification, regression)
 - Modern optimization methods achieve $O(\sqrt{T})$ regret for a general class, and $O(\log T)$ for some special cases

POLA Algorithm

- Consider the following convex sets

$$C_t = \{(A, b) \mid \ell_t(A, b) = 0\}$$

$$C_a = \{(A, b) \mid A \succeq 0, b \geq 1\}$$

- Consider orthogonal projections:

$$P_C(\mathbf{x}) = \operatorname{argmin}_{\mathbf{y} \in C} \|\mathbf{x} - \mathbf{y}\|_2^2$$

- Think of (A, b) as a vector in $d^2 + 1$ dimensional space
- Each step, project onto C_t , then project onto C_a

POLA Algorithm

- Projection onto C_t
 - Let $\mathbf{v}_t = \mathbf{x}_t - \mathbf{x}'_t$
 - Then projection is computed in closed form as

$$\begin{aligned}\alpha_t &= \frac{\ell_t(A_t, b_t)}{\|\mathbf{v}_t\|_2^4 + 1} \\ \hat{A}_t &= A_t - y_t \alpha_t \mathbf{v}_t \mathbf{v}_t^T \\ \hat{b}_t &= b_t + \alpha_t y_t\end{aligned}$$

- Projection onto C_a
 - Orthogonal projection onto positive semi-definite cone obtained by setting negative eigenvalues to 0
 - But, update from A_t to \hat{A}_t was *rank-one*
 - Only 1 negative eigenvalue (interlacing theorem)
 - Thus, only need to compute the smallest eigenvalue and eigenvector (via a power method or related) and subtract it off

Analysis of POLA


- **Theorem:** Let $(\mathbf{x}_1, \mathbf{x}'_1, y_1), \dots, (\mathbf{x}_T, \mathbf{x}'_T, y_T)$ be a sequence of examples and let R be such that $\forall t, R \geq \|\mathbf{x}_t - \mathbf{x}'_t\|_2^4 + 1$. Assume there exists an $A^* \succeq 0$ and $b^* \geq 1$ such that $\ell_t(A^*, b^*) = 0 \forall t$. Then the following bound holds for all $T \geq 1$:

$$\sum_{t=1}^T \ell_t(A_t, b_t)^2 \leq R(\|A^*\|_F^2 + (b^* - b_1)^2).$$

- Note that since $\ell_t(A^*, b^*) = 0$ for all t , the total loss is equal to the regret
- Can generalize this to a regret bound in the case when $\ell_t(A^*, b^*)$ does not always equal 0
- Can also run POLA in batch settings

POLA in Context

- Can we think of POLA in the framework presented earlier?
 - Yes—regularizer is $\|A\|_F^2$ and constraints are defined by the hinge loss ℓ_t
 - Similar to both Schultz and Joachims, and Kwok and Tsang
- We will see later that POLA can also be kernelized to learn non-linear transformations
- In practice, POLA does not appear to be competitive with current state-of-the-art



- Similarly to Schultz and Joachims, utilize relative distance constraints
 - A constraint $(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R}$ has the property that \mathbf{x}_i and \mathbf{x}_j are neighbors of the same class, and \mathbf{x}_i and \mathbf{x}_k are of different classes

[Weinberger, Blitzer, and Saul; NIPS 2005]

- Problem Formulation

- Also define set \mathcal{S} of pairs of points $(\mathbf{x}_i, \mathbf{x}_j)$ such that \mathbf{x}_i and \mathbf{x}_j are neighbors in the same class
- Want to minimize sum of distances of pairs of points in \mathcal{S}
- Also want to satisfy the relative distance constraints

- Mathematically:

$$\begin{aligned} \min_A \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

- Problem Formulation

- Also define set \mathcal{S} of pairs of points $(\mathbf{x}_i, \mathbf{x}_j)$ such that \mathbf{x}_i and \mathbf{x}_j are neighbors in the same class
- Want to minimize sum of distances of pairs of points in \mathcal{S}
- Also want to satisfy the relative distance constraints

- Mathematically:

$$\begin{aligned} \min_{A, \xi} \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_A(\mathbf{x}_i, \mathbf{x}_j) + \gamma \sum_{(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R}} \xi_{ijk} \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & A \succeq 0, \xi_{ijl} \geq 0. \end{aligned}$$

- Introduce slack variables

Comments on LMNN

- Algorithm
 - Special-purpose solver
 - Relies on subgradient computations
 - Ignores inactive constraints
 - Example: MNIST—3.2 billion constraints in 4 hours
 - Software available
- Performance
 - One of the best-performing methods
 - Works in a variety of settings

Learning with Multiple Local Metrics

- Learn several local Mahalanobis metrics instead a single global one
 - Cluster the training data into k partitions
 - Denote c_i as the corresponding cluster for \mathbf{x}_i
 - Learn k Mahalanobis distances A_1, \dots, A_k
- Formulation

$$\begin{aligned} \min_A \quad & \sum_{(\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{S}} d_{A_{c_j}}(\mathbf{x}_i, \mathbf{x}_j) \\ \text{s.t.} \quad & d_{A_{c_k}}(\mathbf{x}_i, \mathbf{x}_k) - d_{A_{c_j}}(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & A_i \succeq 0 \quad \forall i. \end{aligned}$$

- Introduce slack variables as with standard LMNN

[Weinberger and Saul; ICML 2008]

LMNN Results

	Error in %	mnist	20news	letters	isolet	yalefaces
test	LMNN	1.72	14.91	3.62	3.59	6.48
	Multiple Metrics	1.18	13.66	3.2	3.08	6.4
train	LMNN	1.19	9.73	3.54	0.7	3.54
	Multiple Metrics	0.04	7.08	1.55	0	3.57

- Results show improvements using multiple metrics
- Weinberger and Saul also extend LMNN to use ball trees for fast search
 - No time to go into details, see paper

ITML and the LogDet Divergence

- We take the regularizer to be the Log-Determinant Divergence:

$$D_{\ell d}(A, A_0) = \text{trace}(AA_0^{-1}) - \log \det(AA_0^{-1}) - d$$

- Problem formulation:

$$\begin{aligned} \min_A \quad & D_{\ell d}(A, A_0) \\ \text{s.t.} \quad & (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \leq u \quad \text{if } (i, j) \in \mathcal{S} \text{ [similarity constraints]} \\ & (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \geq \ell \quad \text{if } (i, j) \in \mathcal{D} \text{ [dissimilarity constraints]} \end{aligned}$$

[Davis, Kulis, Jain, Sra, and Dhillon; ICML 2007]

LogDet Divergence: Properties

$$D_{\ell d}(A, A_0) = \text{trace}(AA_0^{-1}) - \log \det(AA_0^{-1}) - d,$$

- Properties:

- Scale-invariance

$$D_{\ell d}(A, A_0) = D_{\ell d}(\alpha A, \alpha A_0), \quad \alpha > 0$$

- In fact, for any invertible M

$$D_{\ell d}(A, A_0) = D_{\ell d}(M^T A M, M^T A_0 M)$$

- Expansion in terms of eigenvalues and eigenvectors
($A = V \Lambda V^T$, $A_0 = U \Theta U^T$):

$$D_{\ell d}(A, A_0) = \sum_{i,j} (\mathbf{v}_i^T \mathbf{u}_j)^2 \left(\frac{\lambda_i}{\theta_j} - \log \frac{\lambda_i}{\theta_j} \right) - d$$

Existing Uses of LogDet

- Information Theory
 - Differential relative entropy between two same-mean multivariate Gaussians equal to LogDet divergence between covariance matrices
- Statistics
 - LogDet divergence is known as Stein's loss in the statistics community
- Optimization
 - BFGS update can be written as:

$$\begin{aligned} \min_B \quad & D_{\ell d}(B, B_t) \\ \text{subject to} \quad & B s_t = y_t \quad (\text{"Secant Equation"}) \end{aligned}$$

- $s_t = x_{t+1} - x_t$, $y_t = \nabla f_{t+1} - \nabla f_t$

Key Advantages

- Simple algorithm, easy to implement in Matlab
- Method can be kernelized
- Scales to millions of data points
- Scales to high-dimensional data (text, images, etc.)
- Can incorporate locality-sensitive hashing for *sub-linear time* similarity searches

The Metric Learning Problem

$$D_{\ell d}(A, A_0) = \text{trace}(AA_0^{-1}) - \log \det(AA_0^{-1}) - d$$

- **ITML Goal:**

$$\begin{aligned} \min_A \quad & D_{\ell d}(A, A_0) \\ \text{s.t.} \quad & (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \leq u \quad \text{if } (i, j) \in \mathcal{S} \text{ [similarity constraints]} \\ & (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \geq \ell \quad \text{if } (i, j) \in \mathcal{D} \text{ [dissimilarity constraints]} \end{aligned}$$

Algorithm: Successive Projections

- Algorithm: project successively onto each linear constraint — **converges to globally optimal solution**
- Use projections to update the Mahalanobis matrix:

$$\begin{aligned} \min_A \quad & D_{\ell d}(A, A_t) \\ \text{s.t.} \quad & (\mathbf{x}_i - \mathbf{x}_j)^T A (\mathbf{x}_i - \mathbf{x}_j) \leq u \end{aligned}$$

- Can be solved by $O(d^2)$ rank-one update:

$$A_{t+1} = A_t + \beta_t A_t (\mathbf{x}_i - \mathbf{x}_j) (\mathbf{x}_i - \mathbf{x}_j)^T A_t$$

- Advantages:
 - Automatic enforcement of positive semidefiniteness
 - Simple, closed-form projections
 - No eigenvector calculation
 - Easy to incorporate slack for each constraint

Recent work in Mahalanobis methods

- Recent work has looked at other regularizers, such as $\text{tr}(A)$, which learns low-rank matrices
- Improvements in online metric learning (tighter bounds)
- Kernelization for non-linear metric learning, the topic of the next section

- Online bounds proven for a variant of POLA based on LogDet regularization
 - Combines the best of both worlds
- Minimize the following function at each timestep

$$f_t(A) = D_{\ell d}(A, A_t) + \eta_t \ell_t(A, \mathbf{x}_t, \mathbf{y}_t)$$

- A_t is the current Mahalanobis matrix
- η_t is the learning rate
- $\ell_t(A, \mathbf{x}_t, \mathbf{y}_t)$ is a loss function, e.g.

$$\ell_t(A, \mathbf{x}_t, \mathbf{y}_t) = \frac{1}{2}(d_A(\mathbf{x}_t, \mathbf{y}_t) - p)^2$$

- For appropriate choice of step size, can guarantee $O(\sqrt{T})$ regret
- Empirically outperforms POLA significantly in practice

Non-Mahalanobis methods: Local distance functions

- General approach
 - Learn a distance function for *every* training data point
 - Given m features per point, denote d_m^{ij} as the distance between the m -th feature in points \mathbf{x}_i and \mathbf{x}_j
 - Denote w_m^j as a weight for feature m of point \mathbf{x}_j
 - Then the distance between an arbitrary (e.g., test) image \mathbf{x}_i and a training image \mathbf{x}_j is

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sum_{m=1}^M w_m^j d_m^{ij}$$

- At test time
 - Given test image \mathbf{x}_i , compute $d(\mathbf{x}_i, \mathbf{x}_j)$ between \mathbf{x}_i and every training point \mathbf{x}_j
 - Sort distances to find nearest neighbors

[Frome, Singer, Sha, and Malik; ICCV 2007]

Non-Mahalanobis methods: Local distance functions

Optimization framework

- Denote \mathbf{w}_j as the vector of weights w_m^j
- As before, construct triples (i, j, k) of points such that the distance between \mathbf{x}_i and \mathbf{x}_j should be smaller than the distance between \mathbf{x}_i and \mathbf{x}_k
- Formulate the following problem:

$$\begin{aligned} \min_{\mathbf{W}} \quad & \sum_j \|\mathbf{w}_j\|_2^2 \\ \text{s.t.} \quad & d(\mathbf{x}_i, \mathbf{x}_k) - d(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & \mathbf{w}_j \geq 0 \quad \forall j. \end{aligned}$$

Non-Mahalanobis methods: Local distance functions

Optimization framework

- Denote \mathbf{w}_j as the vector of weights w_m^j
- As before, construct triples (i, j, k) of points such that the distance between \mathbf{x}_i and \mathbf{x}_j should be smaller than the distance between \mathbf{x}_i and \mathbf{x}_k
- Formulate the following problem:

$$\begin{aligned} \min_{\mathbf{W}} \quad & \sum_j \|\mathbf{w}_j\|_2^2 + \gamma \sum_{(i,j,k)} \xi_{ijk} \\ \text{s.t.} \quad & d(\mathbf{x}_i, \mathbf{x}_k) - d(\mathbf{x}_i, \mathbf{x}_j) \geq 1 - \xi_{ijk} \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & \mathbf{w}_j \geq 0 \quad \forall j. \end{aligned}$$

- Introduce slack variables as before
- Very similar to LMNN and other relative distance methods!

Non-Mahalanobis methods: Local distance functions


- Schultz and Joachims

$$\begin{aligned} \min_A \quad & \|A\|_F^2 \\ \text{s.t.} \quad & d_A(\mathbf{x}_i, \mathbf{x}_k) - d_A(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (i, j, k) \in \mathcal{R} \\ & A \succeq 0. \end{aligned}$$

- Frome et al.

$$\begin{aligned} \min_W \quad & \sum_j \|\mathbf{w}_j\|_2^2 \\ \text{s.t.} \quad & d(\mathbf{x}_i, \mathbf{x}_k) - d(\mathbf{x}_i, \mathbf{x}_j) \geq 1 \quad \forall (\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k) \in \mathcal{R} \\ & \mathbf{w}_j \geq 0 \quad \forall j. \end{aligned}$$

Linear Separability



- No linear transformation for this grouping

Kernel Methods



- Map input data to higher-dimensional “feature” space:

$$\mathbf{x} \rightarrow \varphi(\mathbf{x})$$

- Idea: Run machine learning algorithm in feature space
- Use the following mapping:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

Mapping to Feature Space



Kernel Methods

- Map input data to higher-dimensional “feature” space:

$$\mathbf{x} \rightarrow \varphi(\mathbf{x})$$

- Idea: Run machine learning algorithm in feature space
- Use the following mapping:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$$

- Kernel function: $\kappa(\mathbf{x}, \mathbf{y}) = \langle \varphi(\mathbf{x}), \varphi(\mathbf{y}) \rangle$
- “Kernel trick” — no need to explicitly form high-dimensional features
- In this example: $\langle \varphi(\mathbf{x}), \varphi(\mathbf{y}) \rangle = (\mathbf{x}^T \mathbf{y})^2$

Kernel Methods: Short Intro

- Main idea
 - Take an existing learning algorithm
 - Write it using inner products
 - Replace inner products $\mathbf{x}^T \mathbf{y}$ with kernel functions $\varphi(\mathbf{x})^T \varphi(\mathbf{y})$
 - If $\varphi(\mathbf{x})$ is a non-linear function, then algorithm has been *implicitly* non-linearly mapped
- Examples of kernel functions

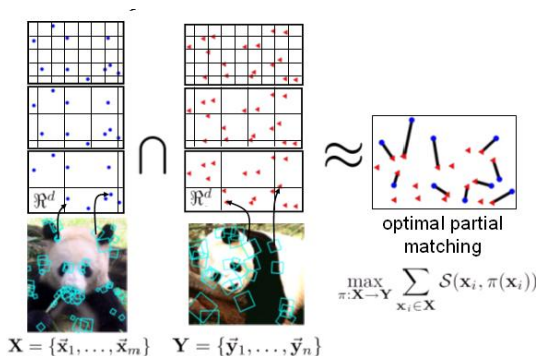
$$\kappa(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y})^p \quad \text{Polynomial Kernel}$$

$$\kappa(\mathbf{x}, \mathbf{y}) = \exp\left(-\frac{\|\mathbf{x} - \mathbf{y}\|_2^2}{2\sigma^2}\right) \quad \text{Gaussian Kernel}$$

$$\kappa(\mathbf{x}, \mathbf{y}) = \tanh(c(\mathbf{x}^T \mathbf{y}) + \theta) \quad \text{Sigmoid Kernel}$$

- Kernel functions also defined over objects such as images, trees, graphs, etc.

Example: Pyramid Match Kernel



- Compute *local* image features
- Perform an approximate matching between features of two images
- Use multi-resolution histograms
- View as a dot product between high-dimensional vectors

[Grauman and Darrell, ICCV 2005]

Example: k -means

Recall the k -means clustering algorithm

- Repeat until convergence:
 - Compute the means of every cluster π_c

$$\mu_c = \frac{1}{|\pi_c|} \sum_{\mathbf{x}_i \in \pi_c} \mathbf{x}_i$$

- Reassign points to their closest mean by computing

$$\|\mathbf{x} - \mu_c\|_2^2$$

for every data point \mathbf{x} and every cluster π_c

Kernelization of k -means

- Expand $\|\mathbf{x} - \mu_c\|_2^2$ as

$$\mathbf{x}^T \mathbf{x} - \frac{2 \sum_{\mathbf{x}_i \in \pi_c} \mathbf{x}^T \mathbf{x}_i}{|\pi_c|} + \frac{\sum_{\mathbf{x}_i, \mathbf{x}_j \in \pi_c} \mathbf{x}_i^T \mathbf{x}_j}{|\pi_c|^2}$$

- No need to explicitly compute the mean; just compute this for every point to every cluster

Example: k -means

Recall the k -means clustering algorithm

- Repeat until convergence:
 - Compute the means of every cluster π_c

$$\mu_c = \frac{1}{|\pi_c|} \sum_{\mathbf{x}_i \in \pi_c} \mathbf{x}_i$$

- Reassign points to their closest mean by computing

$$\|\mathbf{x} - \mu_c\|_2^2$$

for every data point \mathbf{x} and every cluster π_c

Kernelization of k -means

- Expand $\|\mathbf{x} - \mu_c\|_2^2$ as

$$\kappa(\mathbf{x}, \mathbf{x}) - \frac{2 \sum_{\mathbf{x}_i \in \pi_c} \kappa(\mathbf{x}, \mathbf{x}_i)}{|\pi_c|} + \frac{\sum_{\mathbf{x}_i, \mathbf{x}_j \in \pi_c} \kappa(\mathbf{x}_i, \mathbf{x}_j)}{|\pi_c|^2}$$

- Replace inner products with kernels, and this is kernel k -means
- While k -means finds linear separators for the cluster boundaries, kernel k -means finds non-linear separators

Distances vs. Kernel Functions

- Mahalanobis distances:

$$d_A(\mathbf{x}, \mathbf{y}) = (\mathbf{x} - \mathbf{y})^T A (\mathbf{x} - \mathbf{y})$$

- Inner products / kernels:

$$\kappa_A(\mathbf{x}, \mathbf{y}) = \mathbf{x}^T A \mathbf{y}$$

- Algorithms for constructing A learn both measures

From Linear to Nonlinear Learning

Consider the following *kernelized* problem

- You are given a kernel function $\kappa(\mathbf{x}, \mathbf{y}) = \varphi(\mathbf{x})^T \varphi(\mathbf{y})$
- You want to run a metric learning algorithm in kernel space
 - Optimization algorithm cannot use the explicit feature vectors $\varphi(\mathbf{x})$
 - Must be able to compute the distance/kernel over arbitrary points (not just training points)
- Mahalanobis distance is of the form:

$$d_A(\mathbf{x}, \mathbf{y}) = (\varphi(\mathbf{x}) - \varphi(\mathbf{y}))^T A (\varphi(\mathbf{x}) - \varphi(\mathbf{y}))$$

- Kernel is of the form:

$$\kappa_A(\mathbf{x}, \mathbf{y}) = \varphi(\mathbf{x})^T A \varphi(\mathbf{y})$$

- Can be thought of as a kind of *kernel learning* problem

Kernelization of ITML

- First example: ITML
- Recall the update for ITML

$$A_{t+1} = A_t + \beta_t A_t (\mathbf{x}_i - \mathbf{x}_j)(\mathbf{x}_i - \mathbf{x}_j)^T A_t$$

- Distance constraint over pair $(\mathbf{x}_i, \mathbf{x}_j)$
- β_t computed in closed form
- How can we make this update *independent* of the dimensionality?

Kernelization of ITML

- Rewrite the algorithm in terms of inner products (kernel functions)

$$A_{t+1} = A_t + \beta_t A_t (\mathbf{x}_i - \mathbf{x}_j)(\mathbf{x}_i - \mathbf{x}_j)^T A_t$$

- Inner products in this case: $\mathbf{x}_i^T A_t \mathbf{x}_j$

Kernelization of ITML

- Rewrite the algorithm in terms of inner products (kernel functions)

$$X^T A_{t+1} X = X^T A_t X + \beta_t X^T A_t X (\mathbf{e}_i - \mathbf{e}_j) (\mathbf{e}_i - \mathbf{e}_j)^T X^T A_t X$$

- Entry (i, j) of $X^T A_t X$ is exactly $\mathbf{x}_i^T A \mathbf{x}_j = \kappa_A(\mathbf{x}_i, \mathbf{x}_j)$
- Denote $X^T A_t X$ as K_t , the *kernel matrix* at step t

$$K_{t+1} = K_t + \beta_t K_t (\mathbf{e}_i - \mathbf{e}_j) (\mathbf{e}_i - \mathbf{e}_j)^T K_t$$

Kernel Learning

- Squared Euclidean distance in kernel space:

$$\|\mathbf{x}_i - \mathbf{x}_j\|_2^2 = \mathbf{x}_i^T \mathbf{x}_i + \mathbf{x}_j^T \mathbf{x}_j - 2\mathbf{x}_i^T \mathbf{x}_j$$

- Replace with kernel functions / kernel matrix:

$$\kappa(\mathbf{x}_i, \mathbf{x}_i) + \kappa(\mathbf{x}_j, \mathbf{x}_j) - 2\kappa(\mathbf{x}_i, \mathbf{x}_j) = K_{ii} + K_{jj} - 2K_{ij}$$

- Related to ITML, define the following optimization problem

$$\begin{aligned} \min_K \quad & D_{\ell_d}(K, K_0) \\ \text{s.t.} \quad & K_{ii} + K_{jj} - 2K_{ij} \leq u \quad \text{if } (i, j) \in \mathcal{S} \text{ [similarity constraints]} \\ & K_{ii} + K_{jj} - 2K_{ij} \geq \ell \quad \text{if } (i, j) \in \mathcal{D} \text{ [dissimilarity constraints]} \end{aligned}$$

- $K_0 = X^T X$ is the input kernel matrix
- To solve this, only the original kernel function $\kappa(\mathbf{x}_i, \mathbf{x}_j)$ is required

Kernel Learning

- Bregman projections for the kernel learning problem:

$$K_{t+1} = K_t + \beta_t K_t (\mathbf{e}_i - \mathbf{e}_j)(\mathbf{e}_i - \mathbf{e}_j)^T K_t$$

- Suggests a strong connection between the 2 problems
- **Theorem:** Let A^* be the optimal solution to ITML, and $A_0 = I$. Let K^* be the optimal solution to the kernel learning problem. Then $K^* = X^T A^* X$.
 - Solving the kernel learning problem is “equivalent” to solving ITML
 - So we can run entirely in kernel space
 - But, given two new points, how to compute distance?

[Davis, Kulis, Jain, Sra, and Dhillon; ICML 2007]

Induction with LogDet

- **Theorem:** Let A^* be the optimal solution to ITML, and let $A_0 = I$. Let K^* be the optimal solution to the kernel learning problem, and let $K_0 = X^T X$ be the input kernel matrix. Then

$$\begin{aligned}A^* &= I + X S X^T \\ S &= K_0^{-1} (K^* - K_0) K_0^{-1}\end{aligned}$$

- Gives us a way to *implicitly* compute A^* once we solve for K^*
- Algorithm
 - Solve for K^*
 - Construct S using K_0 and K^*
 - Given two points \mathbf{x} and \mathbf{y} , the kernel $\kappa_A(\mathbf{x}, \mathbf{y}) = \mathbf{x}^T A \mathbf{y}$ is computed as

$$\kappa_A(\mathbf{x}_i, \mathbf{x}_j) = \kappa(\mathbf{x}_i, \mathbf{x}_j) + \sum_{i,j=1}^n S_{ij} \kappa(\mathbf{x}, \mathbf{x}_i) \kappa(\mathbf{x}_j, \mathbf{y})$$

Kernelization of POLA

- Recall updates for POLA

$$\begin{aligned}\hat{A}_t &= A_t - y_t \alpha_t \mathbf{v}_t \mathbf{v}_t^T \\ A_{t+1} &= \hat{A}_t - \lambda_d \mathbf{u}_d \mathbf{u}_d^T\end{aligned}$$

- \mathbf{v}_t is the difference of 2 data points
- \mathbf{u}_d is the smallest eigenvector of \hat{A}_t
- 1st update projects onto set C_t where hinge loss is zero (applied only when loss is non-zero)
- 2nd update projects onto PSD cone C_a (applied only when \hat{A}_t has negative eigenvalue)
- Claim:** Analogous to ITML, $A^* = X S X^T$, where X is the matrix of data points
 - Prove this inductively

Kernelization of POLA

Projection onto C_t

- $A_t = XS_tX^T$
- Say the 2 data points are indexed by i and j
 - Then $\mathbf{v}_t = X(\mathbf{e}_i - \mathbf{e}_j)$
- Rewrite $A_t - y_t\alpha\mathbf{v}_t\mathbf{v}_t^T$ to get update from S_t to \hat{S}_t :

$$\begin{aligned}\hat{A}_t &= XS_tX^T - y_t\alpha_tX(\mathbf{e}_i - \mathbf{e}_j)(\mathbf{e}_i - \mathbf{e}_j)^T X^T \\ &= X(S_t - y_t\alpha_t(\mathbf{e}_i - \mathbf{e}_j)(\mathbf{e}_i - \mathbf{e}_j)^T)X^T\end{aligned}$$

Projection onto C_a

- \mathbf{u}_d is an eigenvector of \hat{A}_t , i.e.,

$$\begin{aligned}\hat{A}_t\mathbf{u}_d &= X\hat{S}_tX^T\mathbf{u}_d = \lambda_d\mathbf{u}_d \\ \mathbf{u}_d &= X\left(\frac{1}{\lambda_d}\hat{S}_tX^T\mathbf{u}_d\right) = X\mathbf{q}\end{aligned}$$

- Construction for \mathbf{q} non-trivial; involves kernelized Gram-Schmidt
- Expensive (cubic in dimensionality)

General Kernelization Results

- Recent work by Chatpatanasiri et al. has shown additional kernelization results for
 - LMNN
 - Neighbourhood Component Analysis
 - Discriminant Neighborhood Embedding
- Other recent results show additional, general kernelization results
 - Xing et al.
 - Other regularizers (trace-norm)
- At this point, most/all existing Mahalanobis metric learning methods can be kernelized

- Setup for principal components analysis (PCA)
 - Let $X = [\mathbf{x}_1, \dots, \mathbf{x}_n]$ be a set of data points
 - Typically assume data is centered, not critical here
 - Denote SVD of X as $X = U^T \Sigma V$
 - Left singular vectors in U corresponding to non-zero singular values are an orthonormal basis for the span of the \mathbf{x}_i vectors
 - Covariance matrix is $C = XX^T = U^T \Sigma^T \Sigma U$, kernel matrix is $K = X^T X = V^T \Sigma^T \Sigma V$
- Standard PCA recipe
 - Compute SVD of X
 - Project data onto leading singular vectors U , e.g., $\tilde{\mathbf{x}} = U\mathbf{x}$

Kernel PCA

- Key result from the late 1990s: kernelization of PCA
 - Can also form projections using the kernel matrix
 - Allows one to avoid computing SVD
 - If $X = U^T \Sigma V$, then $U = \Sigma^{-1} V X^T$

$$U \mathbf{x} = \Sigma^{-1} V X^T \mathbf{x}$$

- Computation involves inner products $X^T \mathbf{x}$, eigenvectors V of the kernel matrix, and eigenvalues of the kernel matrix
- Relation to Mahalanobis distance methods
 - Kernel PCA allows one to implicitly compute an orthogonal basis U of the data points, and to project arbitrary data points onto this basis
 - For a data set of n points, dimension of basis is at most n
 - Projecting onto U results in an n -dimensional vector

Using kernel PCA for metric learning

- Given a set of points in kernel space $X = [\varphi(\mathbf{x}_1), \dots, \varphi(\mathbf{x}_n)]$
 - Form a basis U and project data onto that basis to form $\tilde{X} = [\tilde{x}_1, \dots, \tilde{x}_n] = [U\varphi(\mathbf{x}_1), \dots, U\varphi(\mathbf{x}_n)]$ using kernel PCA
- Consider a general unconstrained optimization problem f that is a function of kernel function values, i.e.

$$f(\{\varphi(\mathbf{x}_i)^T A \varphi(\mathbf{x}_j)\}_{i,j=1}^n)$$

- Associated minimization

$$\min_{A \succeq 0} f(\{\varphi(\mathbf{x}_i)^T A \varphi(\mathbf{x}_j)\}_{i,j=1}^n)$$

- Theorem:** The optimal value of the above optimization is the same as that of

$$\min_{A' \succeq 0} f(\{\tilde{x}_i^T A' \tilde{x}_j\}_{i,j=1}^n)$$

where A' is $n \times n$.

Consequences

- Any Mahalanobis distance learning method that is unconstrained and can be expressed as a function of learned inner products can be kernelized
- Examples
 - Neighbourhood Components Analysis
 - LMNN (write as unconstrained via the hinge loss)
 - Discriminant neighborhood embedding
- Generalizing to new points
 - For a new point $\varphi(\mathbf{x})$, construct $\tilde{\mathbf{x}}$ and use Mahalanobis distance with learned matrix A'
- Algorithms
 - Exactly the same algorithms employed as in linear case

Extensions

- Chatpatanasiri et al. considered extensions for low-rank transformations
 - Also showed benefits of kernelization in several scenarios
- Recent results (Jain et al.) have shown complementary results for *constrained* optimization problems
 - ITML is a special case of this analysis
 - Other methods follow easily, e.g., methods based on trace-norm regularization
- Now most Mahalanobis metric learning methods have been shown to be kernelizable

Scalability in Kernel Space

- In many situations, dimensionality d and the number of data points n is high
 - Typically, linear Mahalanobis metric learning methods scale as $O(d^2)$ or $O(d^3)$
 - Kernelized Mahalanobis methods scale as $O(n^2)$ or $O(n^3)$
 - What to do when both are large?
- Main idea: restrict the basis used for learning the metric
 - Can be applied to most methods

Scalability with the kernel PCA approach

- Recall the kernel PCA approach
 - Project onto U , the top n left singular vectors
 - Instead, project onto the top r left singular vectors
 - Proceed as before
- Similar approach can be used for ITML
 - The learned kernel is of the form

$$\kappa_A(\mathbf{x}_i, \mathbf{x}_j) = \kappa(\mathbf{x}_i, \mathbf{x}_j) + \sum_{i,j=1}^n S_{ij} \kappa(\mathbf{x}, \mathbf{x}_i) \kappa(\mathbf{x}_j, \mathbf{y})$$

- Restrict S to be $r \times r$ instead of $n \times n$, where $r < n$ data points are chosen
- Rewrite optimization problem using this form of the kernel
- Constraints on learned distances are still linear, so method can be generalized
- Both approaches can be applied to very large data sets
 - Example: ITML has been applied to data sets of nearly 1 million points (of dimensionality 24,000)

Nearest neighbors with Mahalanobis metrics

- Once metrics are learned, k -nn is typically used
 - k -nn is expensive to compute
 - Must compute distances to all n training points
- Recent methods attempt to speed up NN computation
 - Locality-sensitive hashing
 - Ball trees
- One challenge: can such methods be employed even when algorithms are used in kernel space?
 - Recent work applied in computer vision community has addressed this problem for fast image search

Other non-linear methods

- Recall that kernelized Mahalanobis methods try to learn the distance function

$$\|G\varphi(\mathbf{x}) - G\varphi(\mathbf{y})\|_2^2$$


- Chopra et al. learn the non-linear distance

$$\|G_W(\mathbf{x}) - G_W(\mathbf{y})\|_2^2$$

- G_W is a non-linear function
- Application was face verification
- Algorithmic technique: convolutional networks

[Chopra, Hadsell, and LeCun; CVPR 2005]

Other non-linear methods

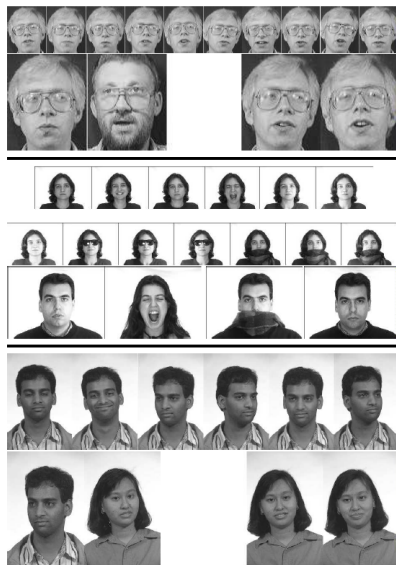


- Setup uses relative distance constraints
 - Denote D_{ij} as the mapped distance between points i and j
 - Let $(\mathbf{x}_i, \mathbf{x}_j, \mathbf{x}_k)$ be a tuple such that $D_{ij} < D_{ik}$ desired
 - The authors define a loss function for each triple of the form

$$Loss = \alpha_1 D_{ij} + \alpha_2 \exp(-\alpha_3 \sqrt{D_{ik}})$$

- Minimize the sum of the losses over all triples
- Metric is trained using a convolutional network with a Siamese architecture from the pixel level

Other non-linear methods





Application: Learning Music Similarity

Comparison of metric learning methods for learning music similarity

- MP3s downloaded from a set of music blogs
 - After pruning: 319 blogs, 164 artists, 74 distinct albums
 - Thousands of songs
- The Echo Nest used to extract features for each song
 - Songs broken up into segments (80ms to a few seconds)
 - Mean segment duration
 - Track tempo estimate
 - Regularity of the beat
 - Estimation of the time signature
 - Overall loudness estimate of the track
 - Estimated overall tatum duration
 - In total, 18 features extracted for each song
- Training done via labels based on blog, artist, and album (separately)

[Slaney, Weinberger, and White; ISMIR 2008]

Application: Learning Music Similarity



Application: Object Recognition

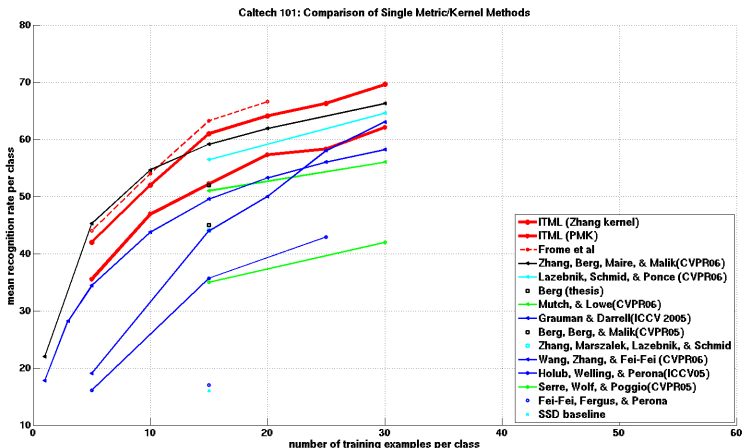


- Several metric learning methods have been evaluated on the Caltech 101 dataset, a benchmark for object recognition set size

Application: Object Recognition

- Used the Caltech-101 data set
 - Standard benchmark for object recognition
 - Many many results for this data set
 - 101 classes, approximately 4000 total images
- Learned metrics over 2 different image embeddings for ITML: pyramid match kernel (PMK) embedding and the embedding from Zhang et al, 2006
- Also learned metrics via Frome et al's local distance function approach
- Computed k -nearest neighbor accuracy over varying training set size and compared to existing results

Application: Object Recognition



Results: Clarify

- Representation: System collects program features during run-time
 - Function counts
 - Call-site counts
 - Counts of program paths
 - Program execution represented as a vector of counts
- Class labels: Program execution errors
- Nearest neighbor software support
 - Match program executions
 - Underlying distance measure should reflect this similarity
- Results
 - LaTeX Benchmark: Error drops from 30% to 15%
 - LogDet is the best performing algorithm across all benchmarks

[Davis, Kulis, Jain, Sra, and Dhillon; ICML 2007]

Application: Human Body Pose Estimation



Pose Estimation



- 500,000 synthetically generated images
- Mean error is 34.5 cm per joint between two random images

Pose Estimation Results

Method	m	$k=1$
L_2 linear scan	24K	8.9
L_2 hashing	24K	9.4
PSH, linear scan	1.5K	9.4
PCA, linear scan	60	13.5
PCA+LogDet, lin. scan	60	13.1
LogDet linear scan	24K	8.4
LogDet hashing	24K	8.8


- Error above given is mean error in cm per joint
- Linear scan requires 433.25 seconds per query; hashing requires 1.39 seconds per query (hashing searches 0.5% of database)

[Jain, Kulis, and Grauman; CVPR 2008]


Pose Estimation Results



Application: Text Retrieval



(a) Classic3



(a) Newsgroups Politics Subset

[Davis and Dhillon; SIGKDD 2008]

Summary and Conclusions

- Metric learning is a mature technology
 - Complaints about scalability in terms of dimensionality or number of data points no longer valid
 - Many different formulations have been studied, especially for Mahalanobis metric learning
 - Online vs offline settings possible
- Metric learning has been applied to many interesting problems
 - Language problems
 - Music similarity
 - Pose estimation
 - Image similarity and search
 - Face verification

Summary and Conclusions

- Metric learning has interesting theoretical components
 - Analysis of online settings
 - Analysis of high-dimensional (kernelized) settings
- Metric learning is still an interesting area of study
 - Learning multiple metrics over data sets
 - New applications
 - Formulations that integrate better with problems other than k -nn
 - Improved algorithms for better scalability
 - ...